

**Assignment - 4**  
**LSTM for Text Classification**

Assignment submission	4 November 2022
Student Name	K.Ragapoorvaja
Student Roll Number	951919CS062
Maximum Marks	2 Marks

**1. Download the Dataset 2. Import required library** import pandas as pd

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import
import LabelEncoder from keras.models
import Model from keras.layers
import LSTM, Activation, Dense, Dropout, Input, Embedding from keras.optimizers
import RMSprop from keras.preprocessing.text
import Tokenizer from
keras.utils
import pad_sequences from keras.utils
import to_categorical from
keras.callbacks
import EarlyStopping
%matplotlib inline
```

**3. Read dataset and do pre-processing** Load the data into Pandas

```
dataframe df =
pd.read_csv('/content/spam.csv',delimiter=',',encoding='latin-1')
df.head()
```

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf. he lives aro...	NaN	NaN	NaN

Drop the columns that are not required for the neural network.

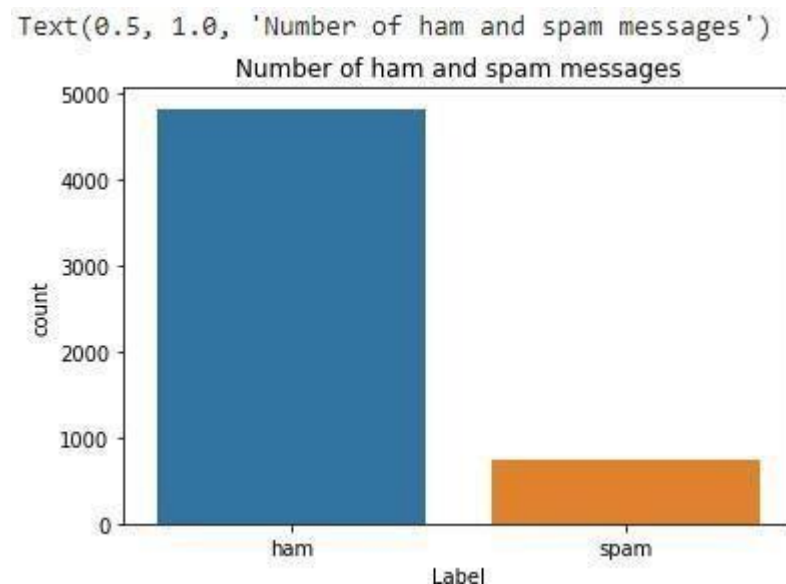
```
df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'],axis=1,inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    v1      5572 non-null     object
1    v2      5572 non-null     object
dtypes: object(2)
memory usage: 87.2+ KB
```

Understand the distribution better.

```
sns.countplot(df.v1) plt.xlabel('Label')
```

```
plt.title('Number of ham and spam messages')
```



- Create input and output vectors. □
- Process the labels.

```
X = df.v2 Y = df.v1 le = LabelEncoder()
```

```
Y = le.fit_transform(Y)
```

```
Y = Y.reshape(-1,1)
```

Split into training and test data.

X\_train,X\_test,Y\_train,Y\_test = train\_test\_split(X,Y,test\_size=0.15) Process the data

- Tokenize the data and convert the text to sequences.
- Add padding to ensure that all the sequences have the same shape.
- There are many ways of taking the \*max\_len\* and here an arbitrary length of 150 is chosen.

```
max_words = 1000 max_len = 150 tok =
```

```
Tokenizer(num_words=max_words)
```

```
tok.fit_on_texts(X_train)
```

```
sequences = tok.texts_to_sequences(X_train)
```

```
sequences_matrix = pad_sequences(sequences,maxlen=max_len)
```

## 5. Create Model

### • Add Layers (LSTM, Dense-(Hidden Layers), Output)

Define the RNN structure. def

RNN():

```
inputs = Input(name='inputs',shape=[max_len]) layer =  
Embedding(max_words,50,input_length=max_len)(inputs) layer =  
LSTM(64)(layer) layer = Dense(256,name='FC1')(layer) layer =  
Activation('relu')(layer) layer = Dropout(0.5)(layer) layer =  
Dense(1,name='out_layer')(layer) layer =  
Activation('sigmoid')(layer) model =  
Model(inputs=inputs,outputs=layer) return model
```

Call the function and compile the model.

```
model = RNN() model.summary()
```

## 6. Compile the Model

```
model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy'])
```

Model: "model"

Layer (type)	Output Shape	Param #
inputs (InputLayer)	[(None, 150)]	0
embedding (Embedding)	(None, 150, 50)	50000
lstm (LSTM)	(None, 64)	29440
FC1 (Dense)	(None, 256)	16640
activation (Activation)	(None, 256)	0
dropout (Dropout)	(None, 256)	0
out_layer (Dense)	(None, 1)	257
activation_1 (Activation)	(None, 1)	0
Total params: 96,337		
Trainable params: 96,337		
Non-trainable params: 0		

## 7. Fit the Model

```
model.fit(sequences_matrix,Y_train,batch_size=128,epochs=10,  
validation_split=0.2,callbacks=[EarlyStopping(monitor='val_loss',min_d  
elta=0.0001)])
```

```
Epoch 1/10  
30/30 [=====] - 11s 286ms/step - loss: 0.3295 - accuracy: 0.8762 - val_loss: 0.1256 - val_accuracy: 0.9757  
Epoch 2/10  
30/30 [=====] - 9s 286ms/step - loss: 0.0880 - accuracy: 0.9797 - val_loss: 0.0440 - val_accuracy: 0.9905  
<keras.callbacks.History at 0x7fadf6edac10>
```

The model performs well on the validation set and this configuration is chosen as the final model.

## 8. Save The Model lstm\_model.save('text\_model.h5')

## 9. Test The Model test\_sequences =

```
tok.texts_to_sequences(X_test) test_sequences_matrix  
=pad_sequences(test_sequences,maxlen=max_len)
```

Evaluate the model on the test set.

```
accr = model.evaluate(test_sequences_matrix,Y_test)
```

```
27/27 [=====] - 1s 23ms/step - loss: 0.0606 - accuracy: 0.9833
```

```
print('Test set\n Loss: {:.3f}\n Accuracy: {:.3f}'.format(accr[0],accr[1]))
```

```
Test set
```

```
Loss: 0.061
```

```
Accuracy: 0.983
```